

Heart Failure Prediction

Gerard Corrales

Gabriel Duffy

Shiley-Marcos School of Engineering, University of San Diego

Abstract

Heart disease remains one of the leading causes of death worldwide, making early and accurate detection crucial for improving patient outcomes. This project investigates the application of machine learning techniques to predict heart disease using a dataset consisting of 918 patient records and 12 clinical features, including variables such as age, blood pressure, cholesterol levels, and other vital health indicators. The primary goal is to assess the predictive accuracy and interpretability of various machine learning models, including CART, C5.0, Random Forest, Naïve Bayes, and Neural Networks.

Extensive data preprocessing was conducted to prepare the dataset for analysis, addressing issues such as outliers, skewed distributions, and the transformation of categorical variables through one-hot encoding. The dataset was split into training and testing sets with a 75/25 ratio, and each model's performance was rigorously evaluated using metrics such as accuracy, precision, recall, and F1-score. Notably, the C5.0 algorithm emerged as the most accurate model, achieving an accuracy rate of 83%. Although the Neural Network model demonstrated a slightly higher accuracy of 87%, it lacked the interpretability that is often crucial in clinical settings.

This project highlights the potential of machine learning in supporting early detection of heart disease, with the C5.0 model offering a compelling balance between accuracy and interpretability. The Random Forest model further provided valuable insights into the importance of specific features, such as ST_Slope and ExerciseAngina, as key predictors. Future work could focus on enhancing predictive performance by exploring ensemble methods or more complex neural network architectures.

Table of Contents

Abstract	2
Introduction	4
Methodology	5
CART	7
C5.0	9
Random Forest	11
Naïve Bayes.....	12
Neural Networks	13
Results	14
Conclusion	15
References	15
Appendix	18

Introduction

Heart disease is one of the leading causes of mortality worldwide (Roth & Mensah, 2020). Early detection and intervention are crucial for reducing heart disease and improving patient results. This project aims to leverage machine learning techniques to predict the presence of heart disease based on a set of descriptive features collected from patients. The first objective of this project is to conduct feature analysis and exploration to identify and understand which features are most indicative of heart disease. This involves exploratory data analysis (EDA) to uncover the relationships and patterns among the variables. Second, to develop and evaluate various classification models to predict heart disease, including CART, C5.0, Random Forests, Naïve Bayes Classification, and Neural Networks. Each model will be evaluated based on criteria such as accuracy, sensitivity, specificity, and misclassification.

The data set used in this project includes a variety of features such as age, resting blood pressure, cholesterol levels, fasting blood sugar, maximum heart rate, exercise-induced angina, and other medical indicators. These features will be preprocessed and standardized to ensure model performance. The expected outcome of this project is to identify the most accurate machine learning model for predicting heart disease.

Methodology

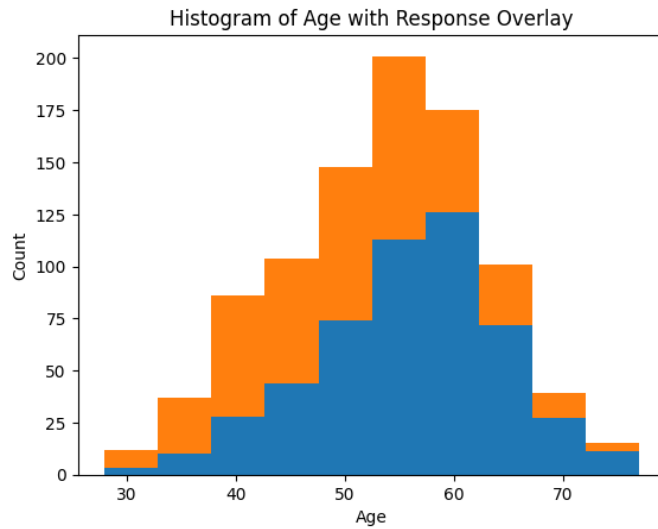
The dataset for this project was obtained from a public repository on GitHub. The dataset contains a total of 918 records and 12 features. The dataset includes a set of features collected from patients, such as age, sex, chest pain, resting blood pressure, cholesterol levels, fasting blood sugar, resting electrocardiogram results, maximum heart rate, exercise-induced angina, numeric value measured in depression (oldpeak), the slope of the peak exercise (ST), and heart disease.

Initial observations revealed no missing values in the dataset. However, outliers were identified in the MaxHR and Oldpeak columns, particularly for the z-scores of MaxHR and Oldpeak, indicating some extreme values that required attention. The dataset also exhibited varying degrees of skewness across features, with notable skewness in the Cholesterol, FastingBS, and Oldpeak columns, among others.


To ensure the data was suitable for analysis, a series of preprocessing steps were carried out. Initially, wrong values in the dataset were handled using K-Nearest Neighbors (KNN) imputation to estimate and fill in the gaps for numerical variables. Categorical variables such as Sex, ChestPainType, RestingECG, ExerciseAngina, and ST_Slope were transformed into numerical format using one-hot encoding, creating dummy variables for each category level. The Oldpeak variable, which had negative values, was standardized using Min-Max scaling to bring all features onto a common scale. This helped in ensuring that all features contribute equally during model training. Identified outliers in the MaxHR and Oldpeak columns were treated to ensure they do not skew the results. Figure 1 and Figure 2 show histograms of age distribution for patients with and without heart disease, with Figure 2 normalizing the distribution to proportionally compare the two groups.

Figure 1

Histogram of Age with Response Overlay

**Figure 2**

Normalized Histogram of Age with Response Overlay



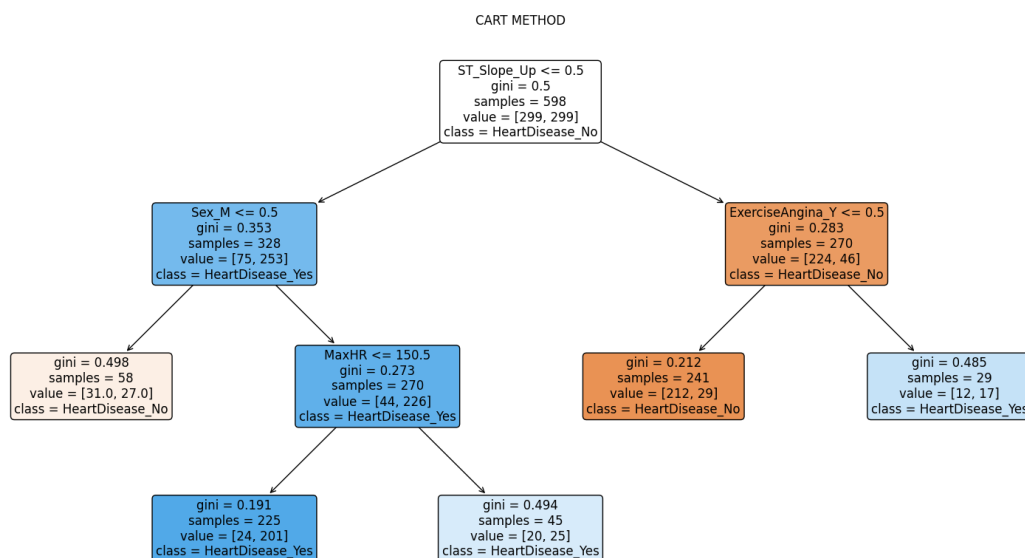
CART

The CART algorithm was employed to develop a predictive model for heart disease. CART is a decision tree algorithm that uses the Gini impurity criterion to split the data into homogenous groups. The CART model was trained on the preprocessed dataset, which included both numerical and categorical variables transformed via one-hot encoding. The dataset was split into training and testing sets, with 75% of the data used for training and 25% reserved for testing. The training set was balanced 50% yes and 50% no responses, it was used to build the decision tree, and it was utilized to evaluate the model's performance.

The model identified several key predictors of heart disease. Among the most significant features were ST_Slope, Sex, MaxHR, and ExerciseAngina (as seen in Figure 3). These variables exhibited the strongest influence on the model's predictions, highlighting their critical role in assessing heart disease risk.

Figure 3

CART Model Decision Tree



The resulting decision tree provided a clear and interpretable set of rules for predicting heart disease. For instance, the root node split on ST_Slope, indicating that patients with a ST_Slope value of less than or equal to 0.5 were more likely to not have heart disease. Subsequent splits on features like Sex and MaxHR further refined the classification, demonstrating that male patients with a MaxHR of less than or equal to 150.5 were more likely to have heart disease. The decision tree also indicated that patients with an ExerciseAngina value of less than or equal to 0.5 were more likely to not have heart disease. The accuracy of the model was found to be 80%, indicating a strong ability to correctly classify patients as having heart disease or not. Additionally, the confusion matrix showed a good balance between sensitivity (true positive rate) and specificity (true negative rate), suggesting that the model effectively identifies both positive and negative cases. In Table 1 shows the results of the classification report.

Table 1*Results of Classification Report*

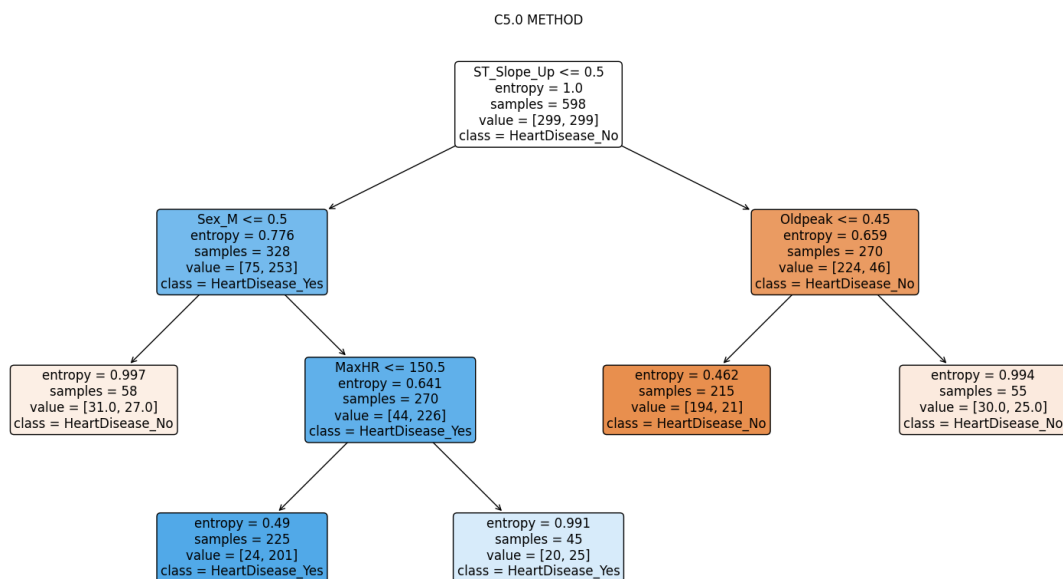
	Precision	Recall	F1-Score	Support
0	0.80	0.77	0.79	111
1	0.80	0.82	0.81	119
Accuracy			0.80	230

C5.0

The C5.0 model was trained on the training dataset. The training and testing splits were identical, with 75% of the data used for training and 25% reserved for testing. The training set was used to construct the decision tree, and the testing set was employed to evaluate the model's performance. The C5.0 model identified several critical predictors of heart disease. The most influential features were ST_Slope, Sex, MaxHR, and Oldpeak. The resulting decision tree provided an interpretable set of rules for heart disease prediction. As seen in Figure 4, the root node split on ST_Slope, indicating that patients with an ST_Slope value of less than or equal to 0.5 were more likely to not have heart disease. For example, male patients ($\text{Sex_M} \leq 0.5$) with a MaxHR of less than or equal to 150.5 were more likely to have heart disease. Additionally, patients with an Oldpeak value of less than or equal to 0.45 were more likely to not have heart disease.

Figure 4

C5.0 Results



The C5.0 model demonstrated robust performance on the test dataset, with an accuracy of 83%. The confusion matrix (Table 2) revealed a good balance between sensitivity and specificity, indicating the model's efficacy in identifying both positive and negative cases of heart disease.

Table 2

C5.0 Confusion Matrix

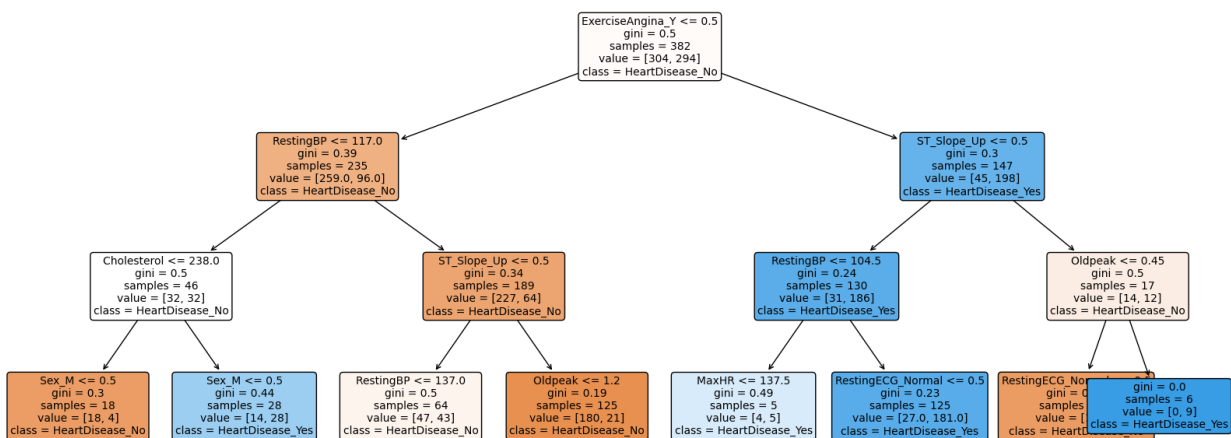
	Precision	Recall	F1-Score	Support
0	0.78	0.90	0.83	111
1	0.89	0.76	0.82	119
Accuracy			0.83	230

Random Forest

The Random Forest model provided valuable insights into the importance of different features in predicting heart disease. As seen in Figure 5, the most important features identified by the model included ExerciseAngina_Y, ST_Slope_Up, Cholesterol and Sex_M, MaxHR and Oldpeak, and RestingBP. ExerciseAngina indicates that patients without exercise-induced angina (ExerciseAngina_Y ≤ 0.5) had a higher likelihood of not having heart disease. The slope of the ST segment during peak exercise was also a crucial factor. Patients with an upward ST slope (ST_Slope_Up ≤ 0.5) showed a lower likelihood of heart disease. Resting blood pressure appeared multiple times in the tree, with various thresholds influencing the likelihood of heart disease. Lower resting blood pressure values were associated with a higher probability of heart disease. Cholesterol levels and gender (Sex_M) also played significant roles in the decision tree, with different thresholds impacting the classification. Maximum heart rate achieved and the ST depression induced by exercise relative to rest (Oldpeak) were other important predictors, influencing subsequent splits in the tree.

Figure 5

Random Forest Decision Tree



Naïve Bayes

When evaluating the Naive Bayes model, the confusion matrix (Table 3) yielded the following results: 88 true negatives (cases where no heart disease was correctly identified), 23 false positives (cases where heart disease was incorrectly predicted), 20 false negatives (cases where heart disease was not predicted but was present), and 99 true positives (cases where heart disease was correctly predicted). This evaluation resulted in an overall model accuracy of 81%. The feature log probabilities for each class (heart disease present or not) were visualized using a heatmap.

Table 3

Naïve Bayes Confusion Matrix


Predicted	False	True	Total
Actual			
0	88	23	111
1	20	99	119
Total	108	122	230

Neural Networks

The model was compiled using the Adam optimizer and binary cross-entropy loss function, suitable for binary classification tasks. It was trained over 150 epochs with a batch size determined automatically by Keras. The training dataset was used to fit the model, achieving an accuracy of approximately 87% (see Figure 6), indicating the model's ability to correctly classify the presence or absence of heart disease. After training, the model was evaluated on the test dataset. While neural networks do not provide feature importance in the same way as decision trees or other interpretable models, the architecture and weights learned during training highlight the complex interactions between features.

Figure 6

Neural Networks Accuracy



Results

The comparative analysis of various machine learning models yielded compelling insights into their effectiveness in predicting heart disease. The C5.0 algorithm demonstrated superior performance, achieving an accuracy of 83%. Its balance between precision and recall was evident in a robust F1-score, underscoring its ability to accurately classify both positive and negative cases. The confusion matrix revealed strong sensitivity and specificity, highlighting the model's reliability in a clinical context where minimizing misclassification is critical.

While the Neural Network model achieved the highest accuracy at 87%, its application was constrained by a lack of interpretability, which poses challenges in healthcare settings requiring transparent decision-making. Despite its predictive strength, the model's opacity in feature contribution limits its practical utility for clinicians. The Random Forest model, with an accuracy of 81%, provided valuable feature importance rankings, identifying variables such as ST_Slope and ExerciseAngina as key predictors of heart disease. This model's capacity to elucidate feature significance is particularly advantageous for understanding the underlying risk factors.

The CART and Naïve Bayes models, achieving accuracies of 80% and 81% respectively, offered strong performance through their simplicity and ease of interpretation. These models provided clear decision rules and probabilistic estimates, making them suitable for scenarios requiring quick, transparent analysis. Overall, the results indicate that while C5.0 offers a well-rounded balance of accuracy and interpretability, Random Forest is particularly beneficial for its insights into feature importance, making both models highly effective for heart disease prediction.

Conclusion

The comprehensive evaluation of machine learning models for heart disease prediction has provided valuable insights into their strengths and practical applications. The C5.0 algorithm, with an accuracy of 83%, proved to be highly effective in balancing precision and recall, making it a robust tool for clinical environments where interpretability is crucial. This model's clear decision rules enable healthcare practitioners to understand and trust the predictions, facilitating informed decision-making. The practical implication of this is that C5.0 can be effectively utilized in diagnostic tools to aid clinicians in identifying patients at risk of heart disease, thus enhancing early intervention and treatment strategies.

Conversely, the Neural Network model achieved the highest accuracy of 87%, demonstrating its superior predictive power. However, its complexity poses challenges for clinical adoption due to the "black box" nature, where the decision-making process is not transparent. Despite its high performance, deploying Neural Networks in clinical settings would necessitate additional efforts to ensure that predictions are interpretable and actionable. This suggests a need for complementary methods or tools that can bridge the gap between high accuracy and practical usability, potentially through explainable AI techniques or hybrid models that combine the strengths of neural networks with more transparent methods.

The Random Forest model, with an accuracy of 81%, provided significant insights into feature importance, such as ST_Slope and ExerciseAngina, which are critical for understanding heart disease risk factors. This feature importance allows for more targeted risk assessments and personalized healthcare strategies, improving the precision of patient evaluations. The CART and Naïve Bayes models, while offering simpler decision rules and probabilistic predictions, can

be valuable for scenarios requiring quick, transparent assessments. Integrating these models into a cohesive strategy, perhaps through ensemble techniques, could enhance overall predictive performance while maintaining clarity and usability. By leveraging the strengths of these various approaches, healthcare providers can develop more effective tools for early heart disease detection, ultimately leading to improved patient outcomes.

References

Larose, C., & Larose, D. (2019). *Data Science Using Python and R*. Wiley.

Tan, P.-N., Steinbach, M., Karpatne, A., & Kumar, V. (2020). *Introduction to data mining* (Second Edition). Pearson.

Roth, G. A., Mensah. (2020). Global Burden of Cardiovascular Diseases and Risk Factors, 1990-2019: Update From the GBD 2019 Study. *Journal of the American College of Cardiology*, 76(25), 2982–3021. <https://doi.org/10.1016/j.jacc.2020.11.010>

Jayachandru001. (2021). *GitHub - jayachandru001/Heart-Failure-Prediction-: This project involves training of Machine Learning models to predict the Heart Failure for Heart Disease event. In this KNN gives a high Accuracy of 89%*. GitHub.

<https://github.com/jayachandru001/Heart-Failure-Prediction-/tree/main>

Appendix

August 9, 2024

0.1 Data Preparation

All The Libraries Used

```
[ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from scipy import stats
import seaborn as sns
from sklearn.model_selection import train_test_split
import random
import statsmodels.tools.tools as stattools
from sklearn.tree import DecisionTreeClassifier, export_graphviz, plot_tree
import graphviz
from sklearn.impute import KNNImputer
from sklearn.metrics import classification_report, confusion_matrix, \
    accuracy_score
from sklearn.ensemble import RandomForestClassifier
from sklearn.naive_bayes import MultinomialNB
from sklearn.preprocessing import MinMaxScaler
import tensorflow as tf
from sklearn import tree
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Input
from sklearn.preprocessing import LabelEncoder
import networkx as nx
```

Importing data set

```
[ ]: heart = pd.read_csv("heart.csv")
```

Visualizing the header

```
[ ]: heart.head()
```

```
[ ]: 
```

	Age	Sex	ChestPainType	RestingBP	Cholesterol	FastingBS	RestingECG	MaxHR	\
0	55	M	NAP	0	0	0	Normal	155	
1	53	M	ASY	80	0	0	Normal	141	
2	32	M	TA	95	0	1	Normal	127	
3	51	M	ASY	95	0	1	Normal	126	
4	57	M	ASY	95	0	1	Normal	182	

	ExerciseAngina	Oldpeak	ST_Slope	HeartDisease
0	N	1.5	Flat	1
1	Y	2.0	Down	0
2	N	0.7	Up	1
3	N	2.2	Flat	1
4	N	0.7	Down	1

Exploring data types

```
[ ]: heart.dtypes
```

```
[ ]: Age                int64
      Sex                object
      ChestPainType      object
      RestingBP          int64
      Cholesterol        int64
      FastingBS          int64
      RestingECG         object
      MaxHR              int64
      ExerciseAngina     object
      Oldpeak            float64
      ST_Slope           object
      HeartDisease       int64
      dtype: object
```

```
[ ]: heart.size
```

```
[ ]: 11016
```

Changing columns object to category type

```
[ ]: for column in ['Sex', 'ChestPainType', 'RestingECG', 'ExerciseAngina', 'ST_Slope']:
      heart[column] = heart[column].astype('category')

      heart.dtypes
```

```
[ ]: Age                int64
      Sex                category
      ChestPainType      category
      RestingBP          int64
      Cholesterol        int64
      FastingBS          int64
      RestingECG         category
      MaxHR              int64
      ExerciseAngina     category
      Oldpeak            float64
      ST_Slope           category
```

```
HeartDisease      int64
dtype: object
```

Checking for missing values

```
[ ]: heart.isna().sum()
```

```
[ ]: Age      0
     Sex      0
     ChestPainType  0
     RestingBP  0
     Cholesterol  0
     FastingBS  0
     RestingECG  0
     MaxHR     0
     ExerciseAngina  0
     Oldpeak   0
     ST_Slope  0
     HeartDisease  0
     dtype: int64
```

Checking for duplicates

```
[ ]: duplicates = heart.duplicated().sum()

     print("Duplicates found: ", duplicates)
```

```
Duplicates found:  0
```

Unique values

```
[ ]: heart.nunique()
```

```
[ ]: Age      50
     Sex      2
     ChestPainType  4
     RestingBP  67
     Cholesterol  222
     FastingBS  2
     RestingECG  3
     MaxHR     119
     ExerciseAngina  2
     Oldpeak   53
     ST_Slope  3
     HeartDisease  2
     dtype: int64
```

Summary of The Central Tendency, Dispersion, and Shape

```
[ ]: heart.describe().T
```

```
[ ]:
count      mean      std   min   25%   50%   75%   max
Age        918.0   53.510893   9.432617  28.0  47.00  54.0  60.0  77.0
RestingBP   918.0  132.396514  18.514154   0.0 120.00 130.0 140.0 200.0
Cholesterol  918.0  198.799564 109.384145   0.0 173.25 223.0 267.0 603.0
FastingBS   918.0   0.233115   0.423046   0.0   0.00   0.0   0.0   1.0
MaxHR       918.0  136.809368  25.460334  60.0 120.00 138.0 156.0 202.0
Oldpeak     918.0   0.887364   1.066570  -2.6   0.00   0.6   1.5   6.2
HeartDisease 918.0   0.553377   0.497414   0.0   0.00   1.0   1.0   1.0
```

Creating an index in our data set

```
[ ]: print("Number of rows: ", heart.shape[0])
      print("Number of columns: ", heart.shape[1])

      # Creating new variable Index
      heart['Index'] = pd.Series(range(0, 918))

      heart.head()
```

Number of rows: 918

Number of columns: 12

```
[ ]:
Age Sex ChestPainType RestingBP Cholesterol FastingBS RestingECG MaxHR \
0  55  M           NAP          0           0           0    Normal    155
1  53  M           ASY          80           0           0    Normal    141
2  32  M           TA           95           0           1    Normal    127
3  51  M           ASY          95           0           1    Normal    126
4  57  M           ASY          95           0           1    Normal    182

ExerciseAngina Oldpeak ST_Slope HeartDisease Index
0              N     1.5     Flat           1      0
1              Y     2.0     Down           0      1
2              N     0.7      Up           1      2
3              N     2.2     Flat           1      3
4              N     0.7     Down           1      4
```

First Visualization to Detect Anomalies in Numerical Data

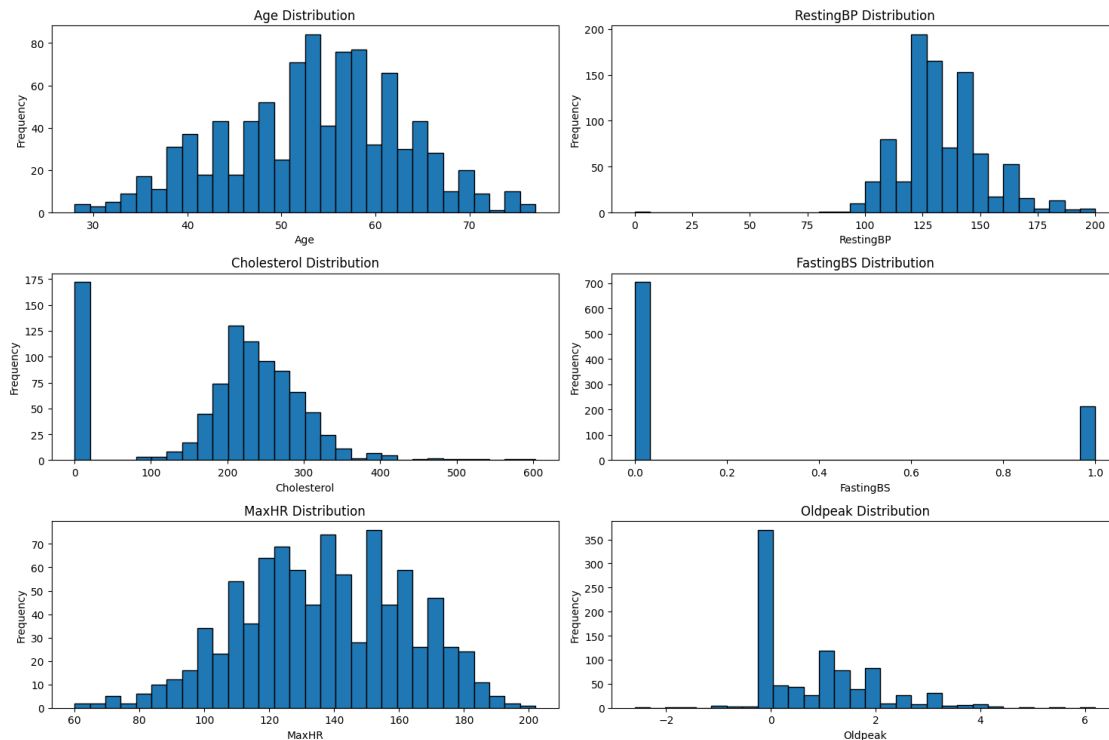
```
[ ]: numerical_columns = ['Age', 'RestingBP', 'Cholesterol', 'FastingBS', 'MaxHR', 'Oldpeak']

fig, axes = plt.subplots(nrows=3, ncols=2, figsize=(15, 10))
axes = axes.flatten()

for i, column in enumerate(numerical_columns):
    axes[i].hist(heart[column], bins=30, edgecolor='black')
    axes[i].set_title(f'{column} Distribution')
    axes[i].set_xlabel(column)
    axes[i].set_ylabel('Frequency')
```

```
plt.tight_layout()
```

```
plt.show()
```



Changing Misleading Field Values Replacing 0 values in “Cholesterol” and “RestingBP” attributes with “nan” because it’s highly unlikely to have 0 cholesterol or 0 Blood pressure.

```
[ ]: heart['Cholesterol'] = heart['Cholesterol'].replace({0: np.nan})
heart['RestingBP'] = heart['RestingBP'].replace({0: np.nan})
```

Showing results after changing misleading values in those columns


```
[ ]: columns_to_show = ['Cholesterol', 'RestingBP']

fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(12, 5))

for i, column in enumerate(columns_to_show):
    axes[i].hist(heart[column], bins=30, edgecolor='black')
    axes[i].set_title(f'{column} Distribution')
    axes[i].set_xlabel(column)
    axes[i].set_ylabel('Frequency')
```

```
plt.tight_layout()
plt.show
```

```
[ ]: <function matplotlib.pyplot.show(close=None, block=None)>
```



Standardizing Numeric Fields to Detect Outliers

```
[ ]: numerical_columns = ['Age', 'RestingBP', 'Cholesterol', 'FastingBS',
                          'MaxHR', 'Oldpeak']
heart_ZScore = pd.DataFrame()

for column in numerical_columns:
    # Calculating Z-score (standardizing)
    heart_ZScore[f'{column}_Z'] = stats.zscore(heart[column])

# Identifying Outliers
heart_age_outliers = heart_ZScore.query('Age_Z > 3 | Age_Z < -3')[['Age_Z']]
heart_resting_outliers = heart_ZScore.query('RestingBP_Z > 3 | RestingBP_Z < -3')[['RestingBP_Z']]
heart_cholesterol_outliers = heart_ZScore.query('Cholesterol_Z > 3 | Cholesterol_Z < -3')[['Cholesterol_Z']]
heart_fasting_outliers = heart_ZScore.query('FastingBS_Z > 3 | FastingBS_Z < -3')[['FastingBS_Z']]
heart_maxhr_outliers = heart_ZScore.query('MaxHR_Z > 3 | MaxHR_Z < -3')[['MaxHR_Z']]
heart_oldpeak_outliers = heart_ZScore.query('Oldpeak_Z > 3 | Oldpeak_Z < -3')[['Oldpeak_Z']]

print("Age")
print(heart_age_outliers)
```



```

print("\nRestingBP_Z")
print(heart_resting_outliers)
print("\nCholesterol_Z")
print(heart_cholesterol_outliers)
print("\nFastingBS_Z")
print(heart_fasting_outliers)
print("\nMaxHR_Z")
print(heart_maxhr_outliers)
print("\nOldpeak_Z")
print(heart_oldpeak_outliers)

```

Age

Empty DataFrame

Columns: [Age_Z]

Index: []

RestingBP_Z

Empty DataFrame

Columns: [RestingBP_Z]

Index: []

Cholesterol_Z

Empty DataFrame

Columns: [Cholesterol_Z]

Index: []

FastingBS_Z

Empty DataFrame

Columns: [FastingBS_Z]

Index: []

MaxHR_Z

MaxHR_Z

126 -3.018469

Oldpeak_Z

Oldpeak_Z

9 -3.271482

208 4.983762

421 4.420905

512 3.858047

707 3.107570

809 3.107570

855 3.295190

Columns “MaxHR_Z” and “Oldpeak_Z” contain many outliers.

Changing target variable values to ‘Yes’ and ‘No’

```
[ ]: Disease_dict = {1: 'Yes', 0: 'No'}

heart['HeartDisease_categorical'] = heart['HeartDisease'].replace(Disease_dict)

heart['HeartDisease_categorical'] = heart['HeartDisease_categorical'].
↳astype('category')

heart['HeartDisease_categorical']
```

```
[ ]: 0      Yes
      1      No
      2      Yes
      3      Yes
      4      Yes
      ...
     913     Yes
     914     Yes
     915     Yes
     916     No
     917     Yes
Name: HeartDisease_categorical, Length: 918, dtype: category
Categories (2, object): ['No', 'Yes']
```

0.2 Exploratory Data Analysis (EDA)

Exploring Categorical Features Using Bar Graph with Response Overlay

```
[ ]: # Creating a contingency table
crosstab_01 = pd.crosstab(heart['Sex'], heart['HeartDisease_categorical'])
crosstab_02 = pd.crosstab(heart['ChestPainType'],
↳heart['HeartDisease_categorical'])
crosstab_03 = pd.crosstab(heart['RestingECG'],
↳heart['HeartDisease_categorical'])
crosstab_04 = pd.crosstab(heart['ExerciseAngina'],
↳heart['HeartDisease_categorical'])
crosstab_05 = pd.crosstab(heart['ST_Slope'], heart['HeartDisease_categorical'])

# Calculating Column Proportions
proportions_01 = round(crosstab_01.div(crosstab_01.sum(0), axis=1) * 100, 1)
proportions_02 = round(crosstab_02.div(crosstab_02.sum(0), axis=1) * 100, 1)
proportions_03 = round(crosstab_03.div(crosstab_03.sum(0), axis=1) * 100, 1)
proportions_04 = round(crosstab_04.div(crosstab_04.sum(0), axis=1) * 100, 1)
proportions_05 = round(crosstab_05.div(crosstab_05.sum(0), axis=1) * 100, 1)

print(proportions_01)
print("\n", proportions_02)
print("\n", proportions_03)
```

```

print("\n", proportions_04)
print("\n", proportions_05)

# Bar graph
crosstab_01.plot(kind='bar', stacked= True)
crosstab_02.plot(kind='bar', stacked= True)
crosstab_03.plot(kind='bar', stacked= True)
crosstab_04.plot(kind='bar', stacked= True)
crosstab_05.plot(kind='bar', stacked= True)

```

HeartDisease_categorical	No	Yes
Sex		
F	34.9	9.8
M	65.1	90.2

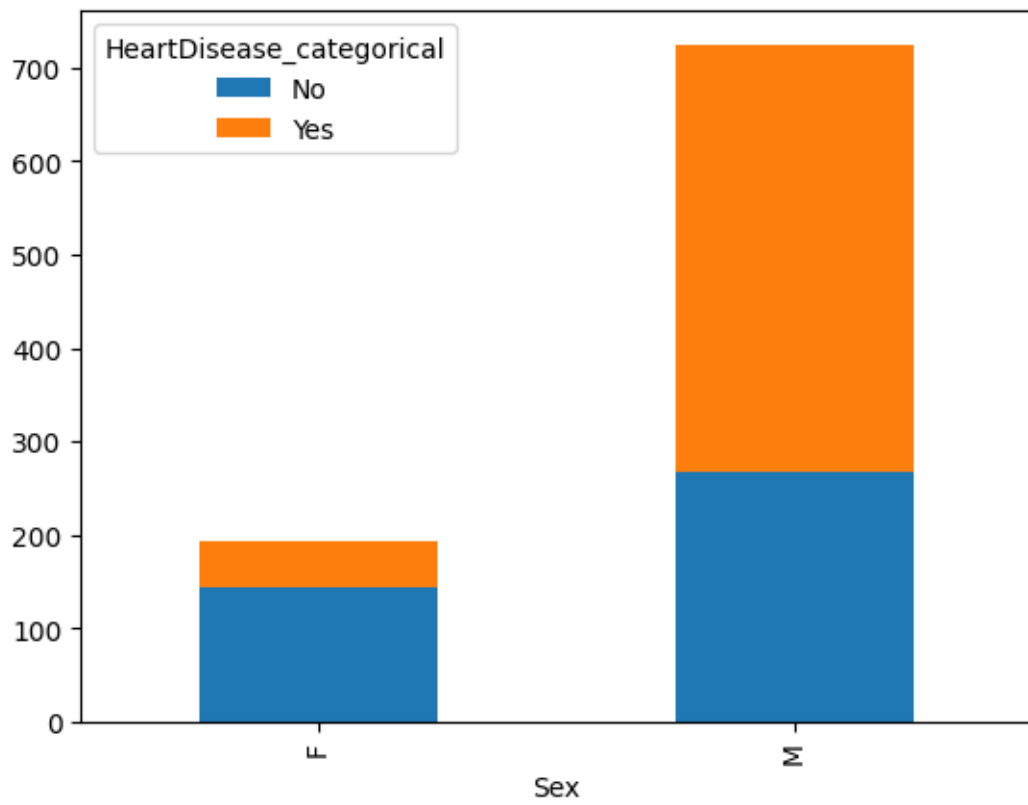
HeartDisease_categorical	No	Yes
ChestPainType		
ASY	25.4	77.2
ATA	36.3	4.7
NAP	32.0	14.2
TA	6.3	3.9


HeartDisease_categorical	No	Yes
RestingECG		
LVH	20.0	20.9
Normal	65.1	56.1
ST	14.9	23.0

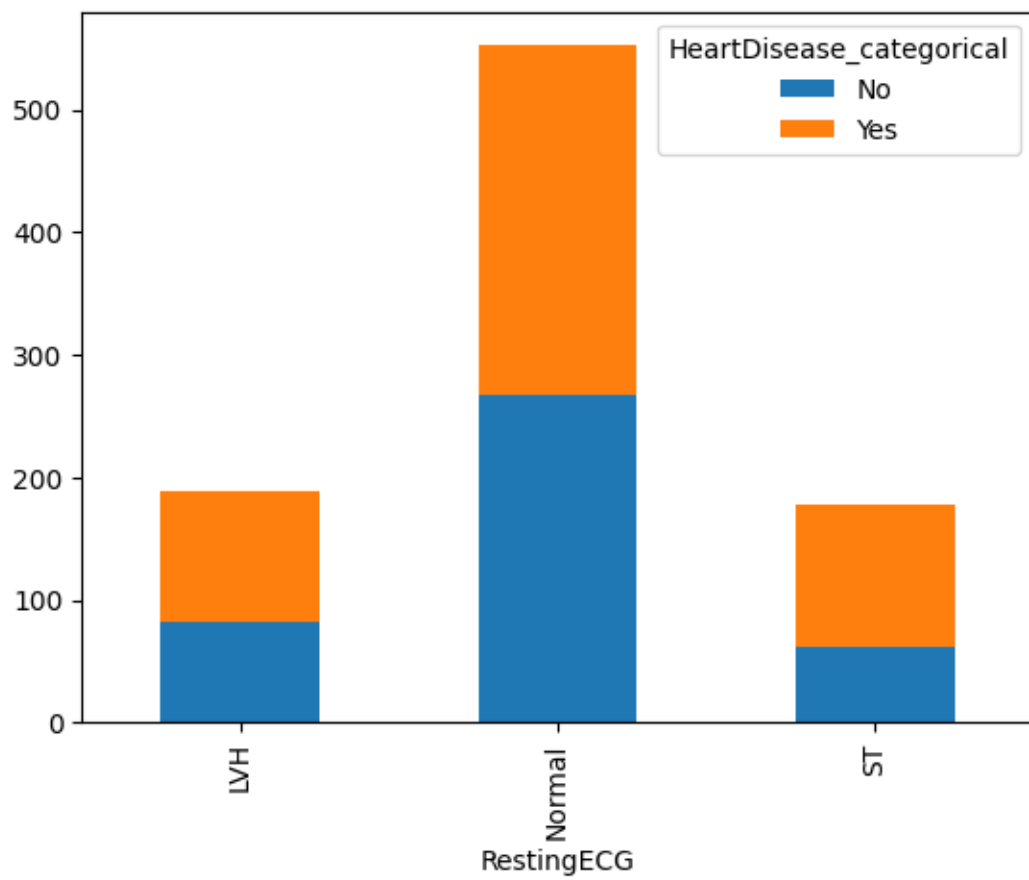
HeartDisease_categorical	No	Yes
ExerciseAngina		
N	86.6	37.8
Y	13.4	62.2

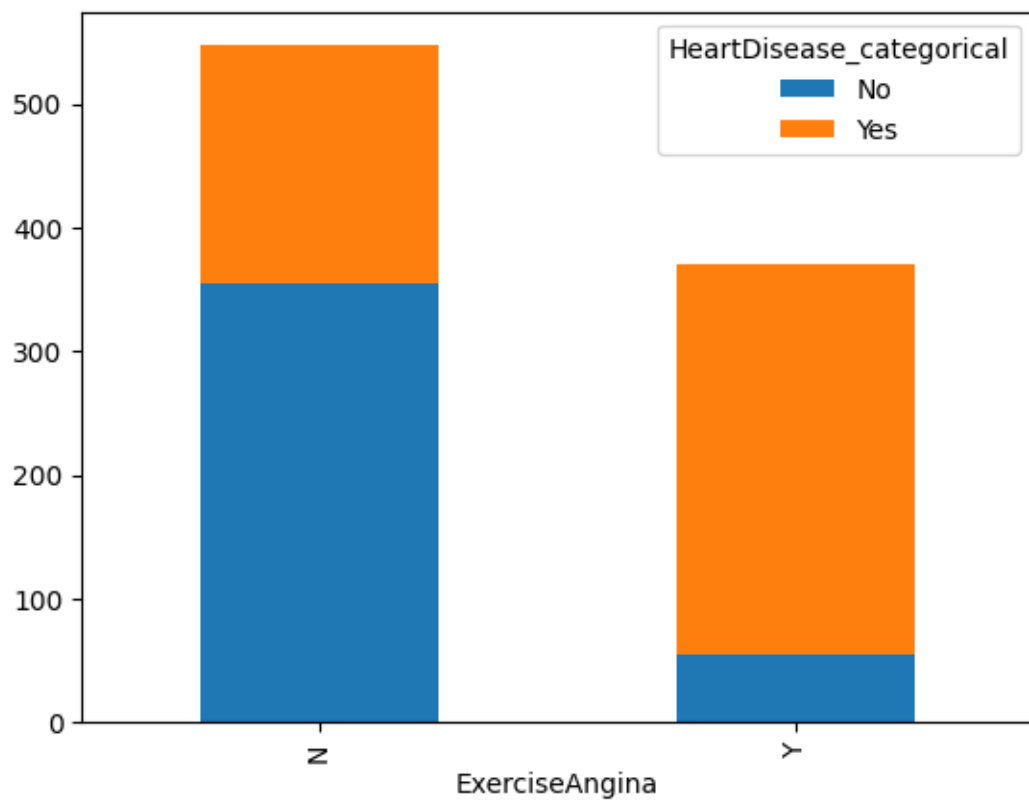
HeartDisease_categorical	No	Yes
ST_Slope		
Down	3.4	9.6
Flat	19.3	75.0
Up	77.3	15.4

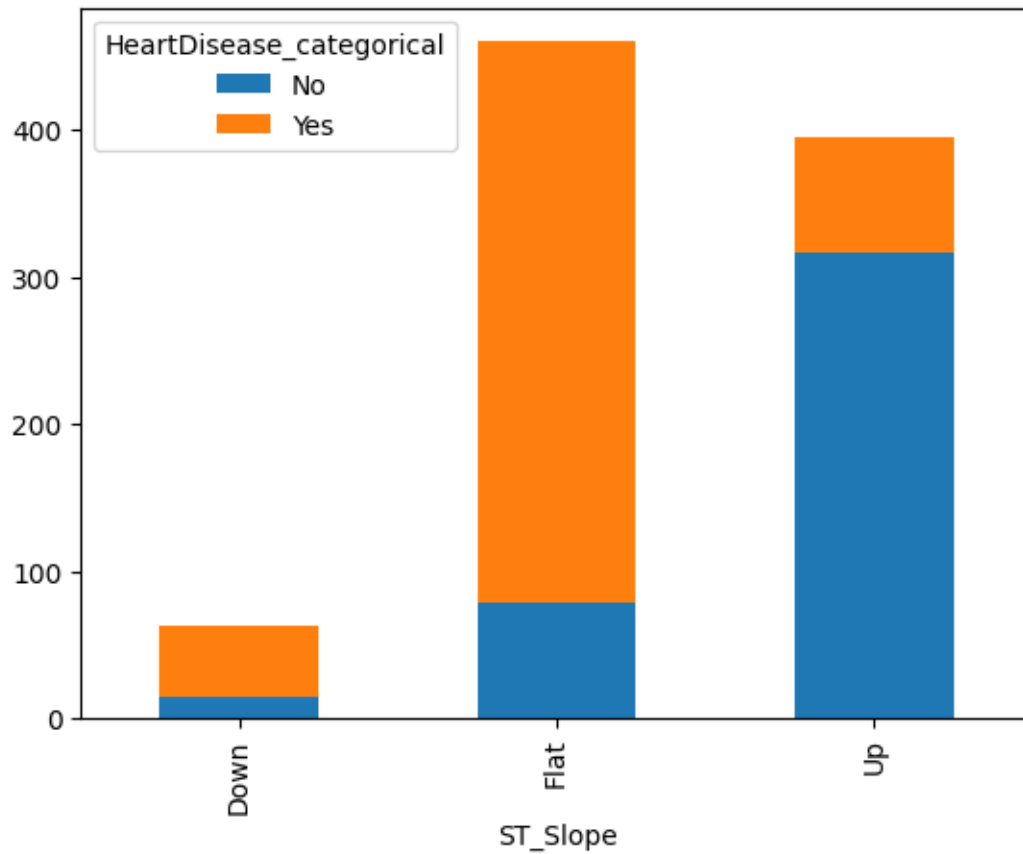
```
[ ]: <Axes: xlabel='ST_Slope'>
```












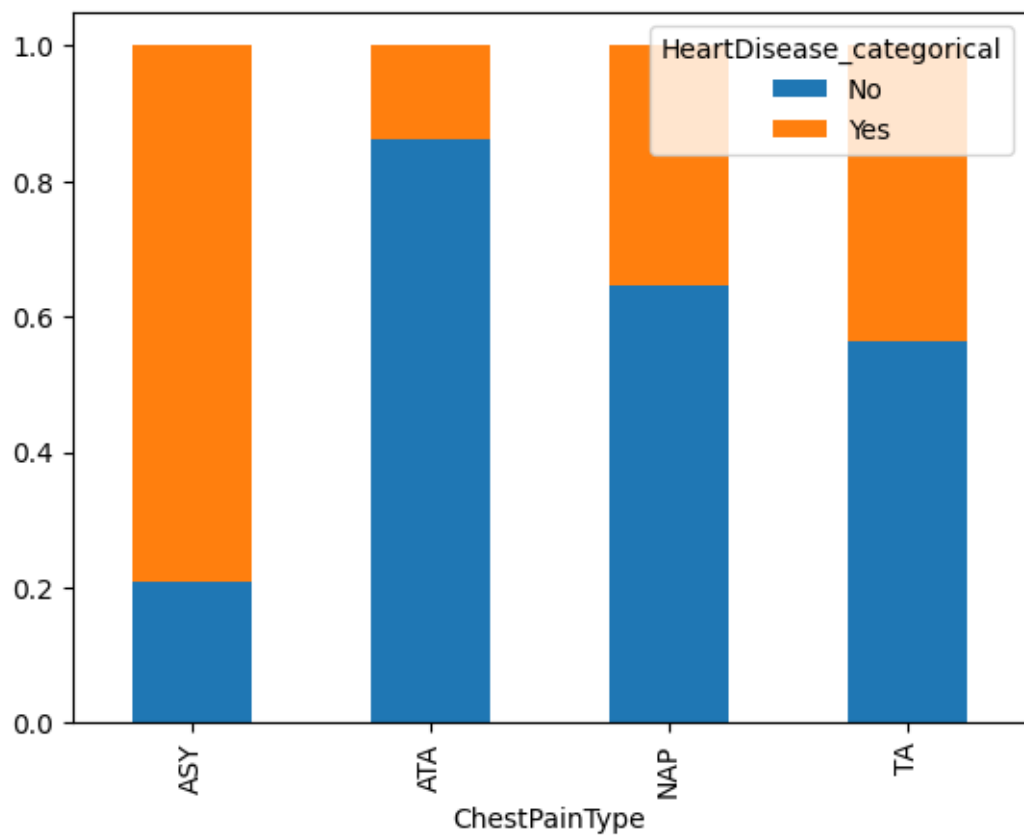
Creating a Normalized Bar Graph

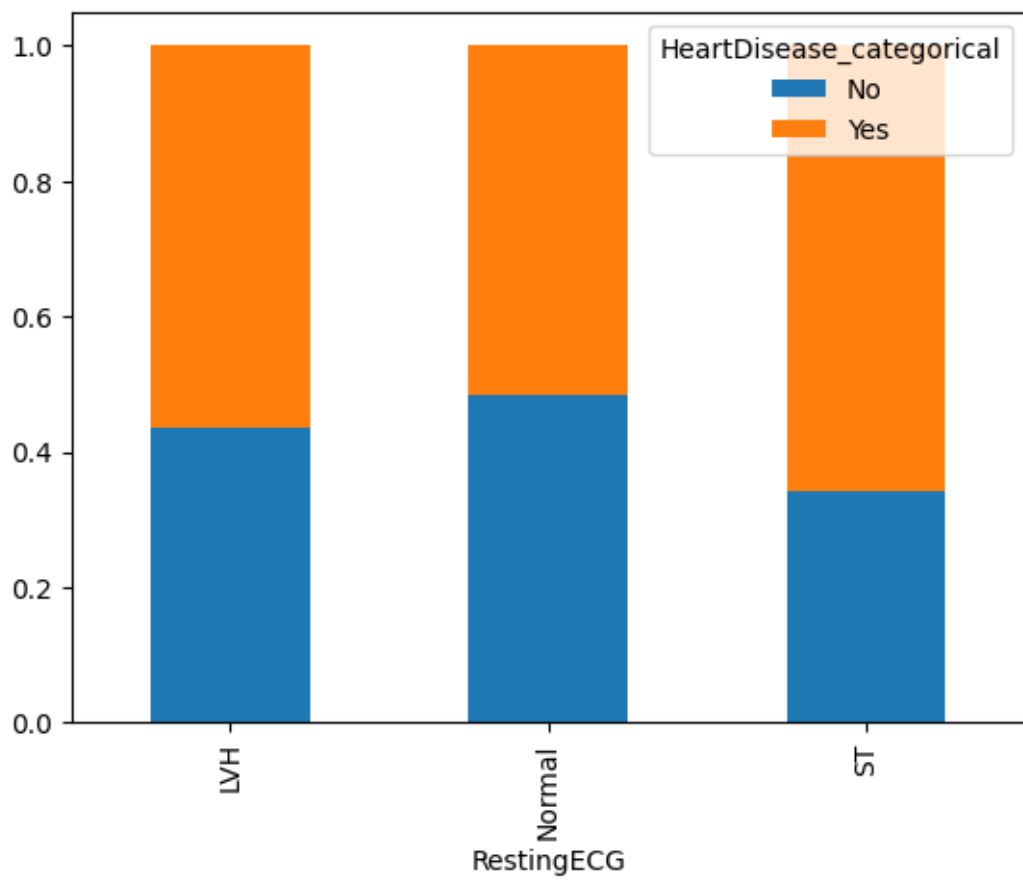
```
[ ]: crosstab_norm_01 = crosstab_01.div(crosstab_01.sum(1), axis= 0)
crosstab_norm_02 = crosstab_02.div(crosstab_02.sum(1), axis= 0)
crosstab_norm_03 = crosstab_03.div(crosstab_03.sum(1), axis= 0)
crosstab_norm_04 = crosstab_04.div(crosstab_04.sum(1), axis= 0)
crosstab_norm_05 = crosstab_05.div(crosstab_05.sum(1), axis= 0)
```

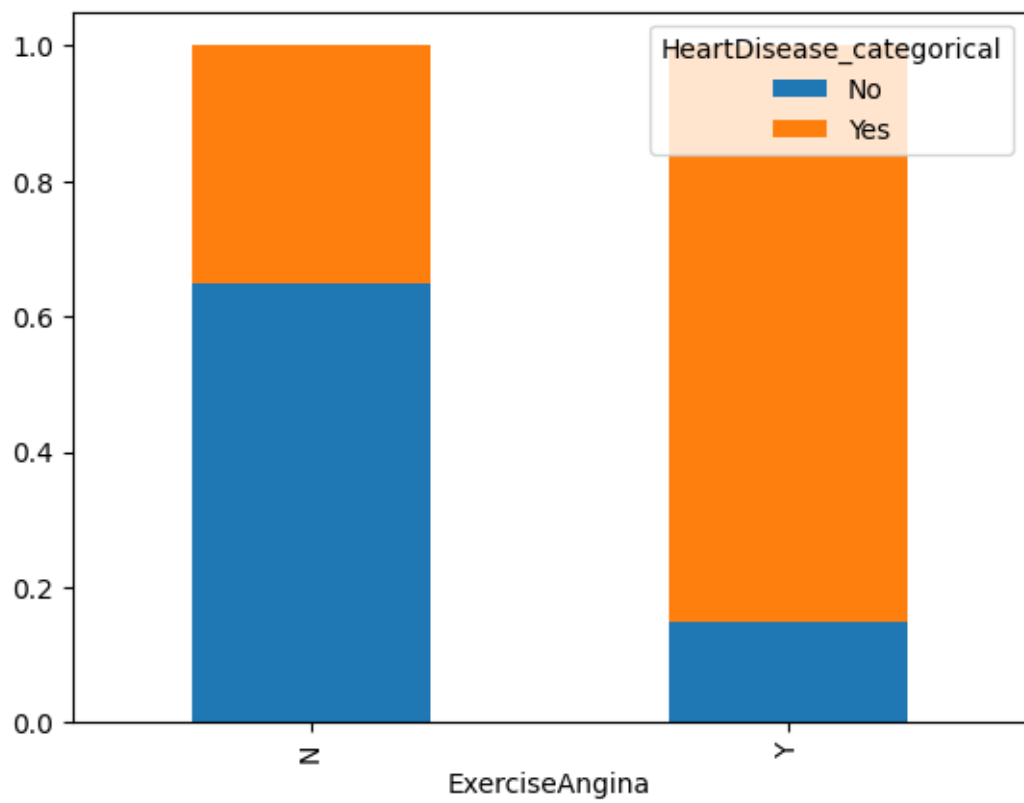
```
crosstab_norm_01.plot(kind='bar', stacked= True)
crosstab_norm_02.plot(kind='bar', stacked= True)
crosstab_norm_03.plot(kind='bar', stacked= True)
crosstab_norm_04.plot(kind='bar', stacked= True)
crosstab_norm_05.plot(kind='bar', stacked= True)
```


```
[ ]: <Axes: xlabel='ST_Slope'>
```









Histogram with Response Overlay for Numerical Data

```
[ ]: #numerical_columns = ['Age', 'RestingBP', 'Cholesterol', 'FastingBS', 'MaxHR', '
    ↳ 'Oldpeak']

heart_age_y = heart[heart.HeartDisease_categorical == "Yes"]['Age']
heart_age_n = heart[heart.HeartDisease_categorical == "No"]['Age']

(n, bins, patches) = plt.hist([heart_age_y, heart_age_n], bins=10, stacked=True)

n_table = np.column_stack((n[0], n[1]))

n_norm = n_table / n_table.sum(axis=1)[:, None]

ourbins = np.column_stack((bins[:10], bins[1:11]))

fig, ax = plt.subplots()


p1 = plt.bar(x=ourbins[:, 0], height=n_norm[:, 0], width=ourbins[:, 1] -
    ↳ ourbins[:, 0], alpha=0.5, edgecolor='black', label='Heart Disease: Yes')
```

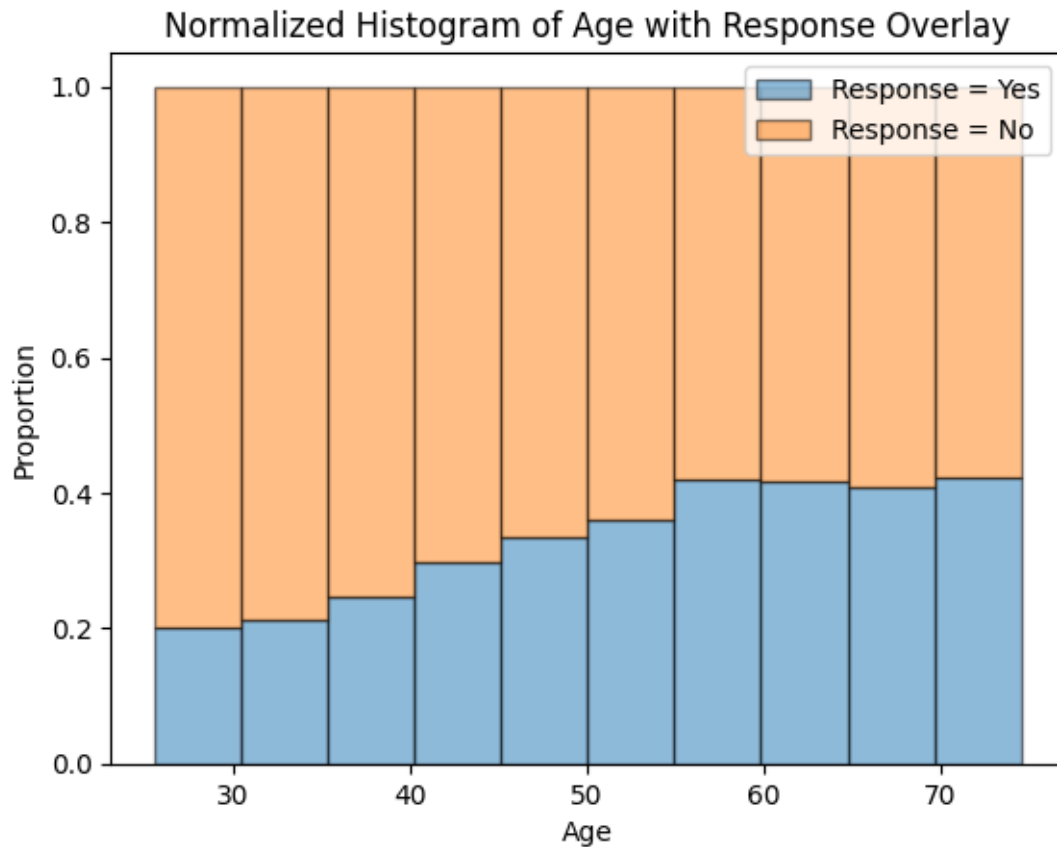
```

p2 = plt.bar(x=ourbins[:, 0], height=n_norm[:, 1], width=ourbins[:, 1] -
↳ourbins[:, 0], bottom=n_norm[:, 0], alpha=0.5, edgecolor='black',
↳label='Heart Disease: No')

plt.legend(['Response = Yes', 'Response = No'])
plt.title('Normalized Histogram of Age with Response Overlay')
plt.xlabel('Age')
plt.ylabel('Proportion')
plt.show()

```

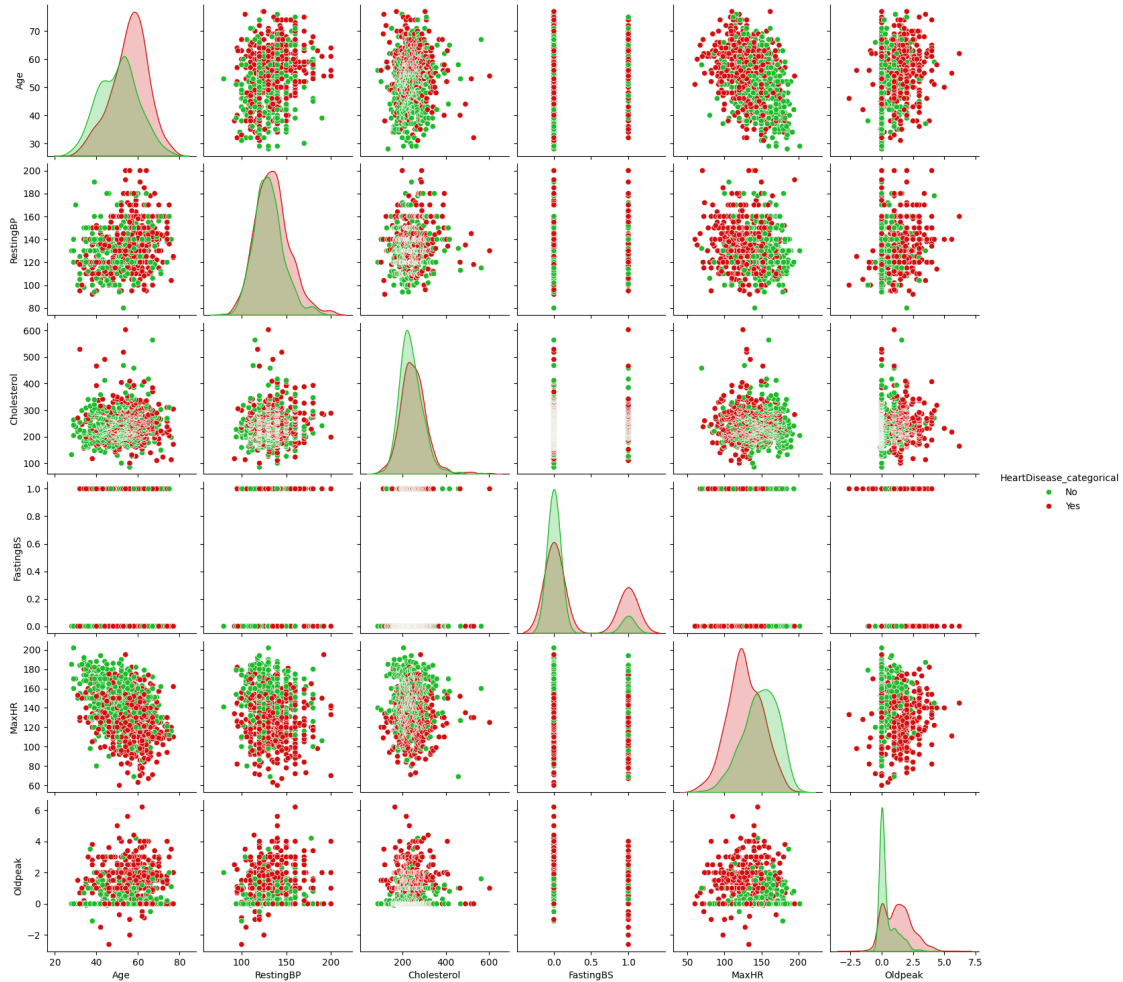




This normalized bar graph shows how the risk of having a heart disease increase with age but at the same time exist a higher number of people without the disease.

```
[ ]: palette = {'No': '#1DBB29', 'Yes': '#D20E0E'}
sns.pairplot(heart,
              vars=['Age', 'RestingBP', 'Cholesterol', 'FastingBS', 'MaxHR', 'Oldpeak'],
              diag_kind="kde",
              hue='HeartDisease_categorical',
              palette=palette)
```

```
[ ]: <seaborn.axisgrid.PairGrid at 0x29d2a9b5c70>
```



0.3 Preparing to Model The Data

One Hot Encoding, Data Imputation, and Partitioning the Data

```
[ ]: # Transforming categorical data into numerical data ( One Hot Encoding)
heart_encoded = pd.get_dummies(heart, columns=['Sex', 'ChestPainType',
                                              'RestingECG',
                                              'ExerciseAngina', 'ST_Slope',
                                              'HeartDisease_categorical'],
                                drop_first=True)

# KNN Imputation
KNN_imputer = KNNImputer(n_neighbors=5)

heart_encoded[['Cholesterol']] = KNN_imputer.
    fit_transform(heart_encoded[['Cholesterol']])
```



```

heart_encoded[['RestingBP']] = KNN_imputer.
↳fit_transform(heart_encoded[['RestingBP']])

# Partitioning the Data
heart_train, heart_test = train_test_split(heart_encoded, test_size= 0.25,↳
↳random_state= 7)

print("Original Data Set:", heart_encoded.shape)
print("Training Data Set:", heart_train.shape, round(heart_train.shape[0] /↳
↳heart_encoded.shape[0] * 100, 2), "%")
print("Test Data Set:", heart_test.shape, round(heart_test.shape[0] /↳
↳heart_encoded.shape[0] * 100, 2), "%")

```

Original Data Set: (918, 18)
Training Data Set: (688, 18) 74.95 %
Test Data Set: (230, 18) 25.05 %

Correlations

```
[ ]: heart_encoded.corr()
```

```
[ ]:
```

	Age	RestingBP	Cholesterol	FastingBS	\
Age	1.000000	0.263081	0.053373	0.198039	
RestingBP	0.263081	1.000000	0.083076	0.067811	
Cholesterol	0.053373	0.083076	1.000000	0.043008	
FastingBS	0.198039	0.067811	0.043008	1.000000	
MaxHR	-0.382045	-0.109662	-0.017239	-0.131438	
Oldpeak	0.258612	0.174220	0.053029	0.052698	
HeartDisease	0.282039	0.117938	0.094071	0.267291	
Index	-0.028882	0.145586	0.672500	-0.198319	
Sex_M	0.055750	0.009425	-0.101706	0.120076	
ChestPainType_ATA	-0.218165	-0.051367	-0.015288	-0.140514	
ChestPainType_NAP	-0.011335	-0.027483	-0.062229	-0.039249	
ChestPainType_TA	0.032042	0.049463	-0.047322	0.026885	
RestingECG_Normal	-0.230566	-0.113718	-0.042407	-0.093028	
RestingECG_ST	0.136798	0.089145	-0.024530	0.127110	
ExerciseAngina_Y	0.215793	0.153008	0.077549	0.060451	
ST_Slope_Flat	0.185568	0.110111	0.093627	0.107006	
ST_Slope_Up	-0.258067	-0.105926	-0.089995	-0.161730	
HeartDisease_categorical_Yes	0.282039	0.117938	0.094071	0.267291	

	MaxHR	Oldpeak	HeartDisease	Index	\
Age	-0.382045	0.258612	0.282039	-0.028882	
RestingBP	-0.109662	0.174220	0.117938	0.145586	
Cholesterol	-0.017239	0.053029	0.094071	0.672500	
FastingBS	-0.131438	0.052698	0.267291	-0.198319	
MaxHR	1.000000	-0.160691	-0.400421	0.175307	
Oldpeak	-0.160691	1.000000	0.403951	0.072980	

HeartDisease	-0.400421	0.403951	1.000000	-0.139166
Index	0.175307	0.072980	-0.139166	1.000000
Sex_M	-0.189186	0.105734	0.305445	-0.181174
ChestPainType_ATA	0.253735	-0.262124	-0.401924	0.123510
ChestPainType_NAP	0.134580	-0.106212	-0.212964	-0.039486
ChestPainType_TA	0.100025	0.032231	-0.054790	0.005106
RestingECG_Normal	0.023801	-0.116719	-0.091580	-0.065735
RestingECG_ST	-0.157879	0.055958	0.102527	-0.105835
ExerciseAngina_Y	-0.370425	0.408752	0.494282	0.024372
ST_Slope_Flat	-0.342581	0.283295	0.554134	-0.004604
ST_Slope_Up	0.383397	-0.450577	-0.622164	0.043002
HeartDisease_categorical_Yes	-0.400421	0.403951	1.000000	-0.139166

	Sex_M	ChestPainType_ATA	ChestPainType_NAP	\
Age	0.055750	-0.218165	-0.011335	
RestingBP	0.009425	-0.051367	-0.027483	
Cholesterol	-0.101706	-0.015288	-0.062229	
FastingBS	0.120076	-0.140514	-0.039249	
MaxHR	-0.189186	0.253735	0.134580	
Oldpeak	0.105734	-0.262124	-0.106212	
HeartDisease	0.305445	-0.401924	-0.212964	
Index	-0.181174	0.123510	-0.039486	
Sex_M	1.000000	-0.161522	-0.066486	
ChestPainType_ATA	-0.161522	1.000000	-0.256767	
ChestPainType_NAP	-0.066486	-0.256767	1.000000	
ChestPainType_TA	-0.004031	-0.110679	-0.122381	
RestingECG_Normal	-0.010634	0.107941	0.005010	
RestingECG_ST	0.063715	-0.046111	-0.042236	
ExerciseAngina_Y	0.190664	-0.300365	-0.166030	
ST_Slope_Flat	0.116077	-0.304667	-0.072031	
ST_Slope_Up	-0.150942	0.357588	0.093583	
HeartDisease_categorical_Yes	0.305445	-0.401924	-0.212964	

	ChestPainType_TA	RestingECG_Normal	\
Age	0.032042	-0.230566	
RestingBP	0.049463	-0.113718	
Cholesterol	-0.047322	-0.042407	
FastingBS	0.026885	-0.093028	
MaxHR	0.100025	0.023801	
Oldpeak	0.032231	-0.116719	
HeartDisease	-0.054790	-0.091580	
Index	0.005106	-0.065735	
Sex_M	-0.004031	-0.010634	
ChestPainType_ATA	-0.110679	0.107941	
ChestPainType_NAP	-0.122381	0.005010	
ChestPainType_TA	1.000000	-0.057719	
RestingECG_Normal	-0.057719	1.000000	

RestingECG_ST	-0.011611	-0.602314
ExerciseAngina_Y	-0.128105	-0.072924
ST_Slope_Flat	-0.010486	-0.047172
ST_Slope_Up	0.002087	0.078563
HeartDisease_categorical_Yes	-0.054790	-0.091580

	RestingECG_ST	ExerciseAngina_Y	ST_Slope_Flat \
Age	0.136798	0.215793	0.185568
RestingBP	0.089145	0.153008	0.110111
Cholesterol	-0.024530	0.077549	0.093627
FastingBS	0.127110	0.060451	0.107006
MaxHR	-0.157879	-0.370425	-0.342581
Oldpeak	0.055958	0.408752	0.283295
HeartDisease	0.102527	0.494282	0.554134
Index	-0.105835	0.024372	-0.004604
Sex_M	0.063715	0.190664	0.116077
ChestPainType_ATA	-0.046111	-0.300365	-0.304667
ChestPainType_NAP	-0.042236	-0.166030	-0.072031
ChestPainType_TA	-0.011611	-0.128105	-0.010486
RestingECG_Normal	-0.602314	-0.072924	-0.047172
RestingECG_ST	1.000000	0.107036	0.043017
ExerciseAngina_Y	0.107036	1.000000	0.382237
ST_Slope_Flat	0.043017	0.382237	1.000000
ST_Slope_Up	-0.058936	-0.455676	-0.870951
HeartDisease_categorical_Yes	0.102527	0.494282	0.554134

	ST_Slope_Up	HeartDisease_categorical_Yes
Age	-0.258067	0.282039
RestingBP	-0.105926	0.117938
Cholesterol	-0.089995	0.094071
FastingBS	-0.161730	0.267291
MaxHR	0.383397	-0.400421
Oldpeak	-0.450577	0.403951
HeartDisease	-0.622164	1.000000
Index	0.043002	-0.139166
Sex_M	-0.150942	0.305445
ChestPainType_ATA	0.357588	-0.401924
ChestPainType_NAP	0.093583	-0.212964
ChestPainType_TA	0.002087	-0.054790
RestingECG_Normal	0.078563	-0.091580
RestingECG_ST	-0.058936	0.102527
ExerciseAngina_Y	-0.455676	0.494282
ST_Slope_Flat	-0.870951	0.554134
ST_Slope_Up	1.000000	-0.622164
HeartDisease_categorical_Yes	-0.622164	1.000000

Skewness of The Distribution

```
[ ]: heart_encoded.skew()
```

```
[ ]: Age                -0.195933
      RestingBP          0.607525
      Cholesterol        1.373396
      FastingBS          1.264484
      MaxHR              -0.144359
      Oldpeak            1.022872
      HeartDisease       -0.215086
      Index              0.000000
      Sex_M              -1.424540
      ChestPainType_ATA   1.595899
      ChestPainType_NAP   1.346107
      ChestPainType_TA    4.130983
      RestingECG_Normal   -0.414489
      RestingECG_ST       1.551033
      ExerciseAngina_Y    0.391329
      ST_Slope_Flat       -0.004364
      ST_Slope_Up         0.282079
      HeartDisease_categorical_Yes -0.215086
      dtype: float64
```

Validating The Data Partition

```
[ ]: from statsmodels.stats.proportion import proportions_ztest

      # Get counts for HeartDisease_categorical in training and test sets
      count_train = np.sum(heart_train['HeartDisease'] == 1)
      count_test = np.sum(heart_test['HeartDisease'] == 1)

      n_train = len(heart_train)
      n_test = len(heart_test)

      # Perform two-sample Z-test for proportion
      count = np.array([count_train, count_test])
      nobs = np.array([n_train, n_test])

      z_stat, p_val = proportions_ztest(count, nobs)

      print(f'Two-sample Z-test for HeartDisease: z-statistic = {z_stat}, p-value = {p_val}')
```

Two-sample Z-test for HeartDisease: z-statistic = 1.268056802569942, p-value = 0.20477766636519124

Based on the Z-test result, there is no significant difference between the training and test sets for the HeartDisease_categorical variable. This indicates that the partitioning of the dataset into training and test sets did not introduce a systematic difference in the proportions of the “Yes” and “No” responses, which is a good indication that the partitioning is valid.

Balancing The Training Data Set

```
[ ]: total_train = heart_train['HeartDisease_categorical_Yes'].value_counts()

Yes_total = total_train.iloc[1] / heart_train.shape[0] * 100
No_total = total_train.iloc[0] / heart_train.shape[0] * 100

print("Total number of 'Yes' in the Training:", round(Yes_total, 2), "%")
print("Total number of 'No' in the Training:", round(No_total, 2), "%")

total_train
```

Total number of 'Yes' in the Training: 43.46 %

Total number of 'No' in the Training: 56.54 %

```
[ ]: HeartDisease_categorical_Yes
True      389
False     299
Name: count, dtype: int64
```

```
[ ]: # Increasing the percentage of "Yes" and "No" responses to 50%
total_train = heart_train['HeartDisease_categorical_Yes'].value_counts()
current_yes_count = total_train[True]
current_no_count = total_train[False]

# Determine the target count for each class to achieve 50-50 balance
target_count = min(current_yes_count, current_no_count)

# Resample to achieve balance
if current_yes_count > target_count:
    # Downsample Yes
    to_downsample_yes = heart_train[heart_train['HeartDisease_categorical_Yes']
    ↪ == True]
    downsampled_yes = to_downsample_yes.sample(n=target_count, random_state=7)
    balanced_train = pd.concat([downsampled_yes,
    ↪ heart_train[heart_train['HeartDisease_categorical_Yes'] == False]])
elif current_no_count > target_count:
    # Downsample No
    to_downsample_no = heart_train[heart_train['HeartDisease_categorical_Yes']
    ↪ == False]
    downsampled_no = to_downsample_no.sample(n=target_count, random_state=7)
    balanced_train = pd.concat([downsampled_no,
    ↪ heart_train[heart_train['HeartDisease_categorical_Yes'] == True]])

# Shuffle the balanced dataset
balanced_train = balanced_train.sample(frac=1, random_state=7).
    ↪ reset_index(drop=True)
```

```
# Verify the balance
balanced_counts = balanced_train['HeartDisease_categorical_Yes'].value_counts()
print("Balanced distribution in the training dataset:")
print(balanced_counts)
```

```
Balanced distribution in the training dataset:
HeartDisease_categorical_Yes
True      299
False     299
Name: count, dtype: int64
```

Modeling Phase, Decision Trees CART Method

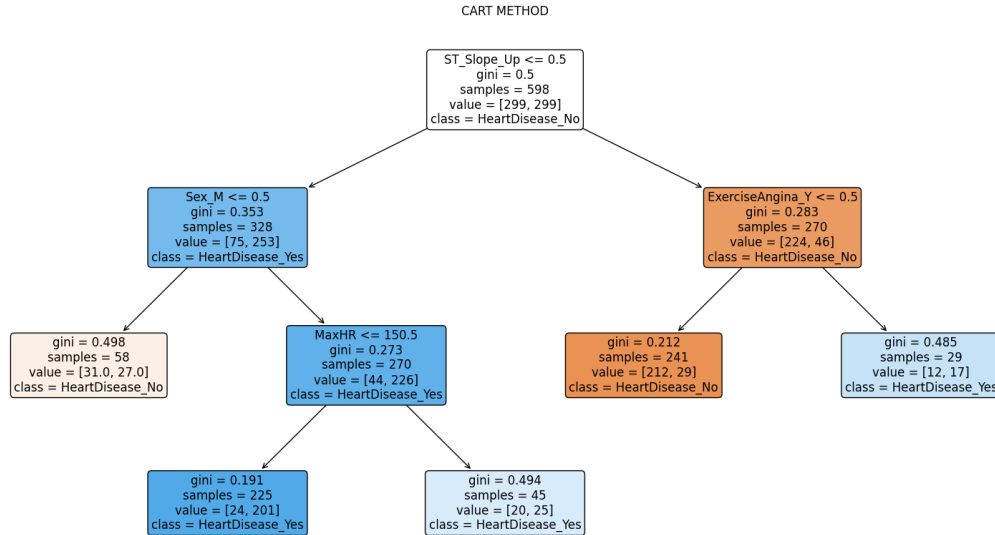
```
[ ]: # Separate data into features and target variable
X = balanced_train.drop(['HeartDisease', 'HeartDisease_categorical_Yes'],
    ↳ 'Index'], axis= 1)
y = balanced_train[['HeartDisease_categorical_Yes']]

# Defining feature names
X_names = list(balanced_train.columns)
X_names.remove('Index')
X_names.remove('HeartDisease')
X_names.remove('HeartDisease_categorical_Yes')

y_names = ["HeartDisease_No", 'HeartDisease_Yes']

# Running the CART algorithm / training
cart01 = DecisionTreeClassifier(criterion= 'gini', max_leaf_nodes= 5).fit(X, y)

# Visualize the decision tree
plt.figure(figsize=(20,10))
plot_tree(cart01, feature_names=X_names, class_names=y_names, filled=True,
    ↳ rounded=True, fontsize=12)
plt.title("CART METHOD")
plt.show()
```



Making predictions

```
[ ]: feature_columns = [
    'Age', 'RestingBP', 'Cholesterol', 'FastingBS', 'MaxHR', 'Oldpeak',
    'Sex_M', 'ChestPainType_ATA', 'ChestPainType_NAP', 'ChestPainType_TA',
    'RestingECG_Normal', 'RestingECG_ST', 'ExerciseAngina_Y', 'ST_Slope_Flat',
    ↪ 'ST_Slope_Up'
]

# Selecting feature columns
X_test = heart_test[feature_columns]

# Making predictions
predHeartDiseaseCART = cart01.predict(X_test)

# Selecting response variable
y_test = heart_test['HeartDisease']

# Results
print("Confusion Matrix:")
print(confusion_matrix(y_test, predHeartDiseaseCART))
print("\nClassification Report:")
print(classification_report(y_test, predHeartDiseaseCART, target_names=["0",
    ↪ "1"]))
print("\nAccuracy Score:")
print(accuracy_score(y_test, predHeartDiseaseCART))
```

Confusion Matrix:

[[86 25]

[21 98]]

Classification Report:

	precision	recall	f1-score	support
0	0.80	0.77	0.79	111
1	0.80	0.82	0.81	119
accuracy			0.80	230
macro avg	0.80	0.80	0.80	230
weighted avg	0.80	0.80	0.80	230

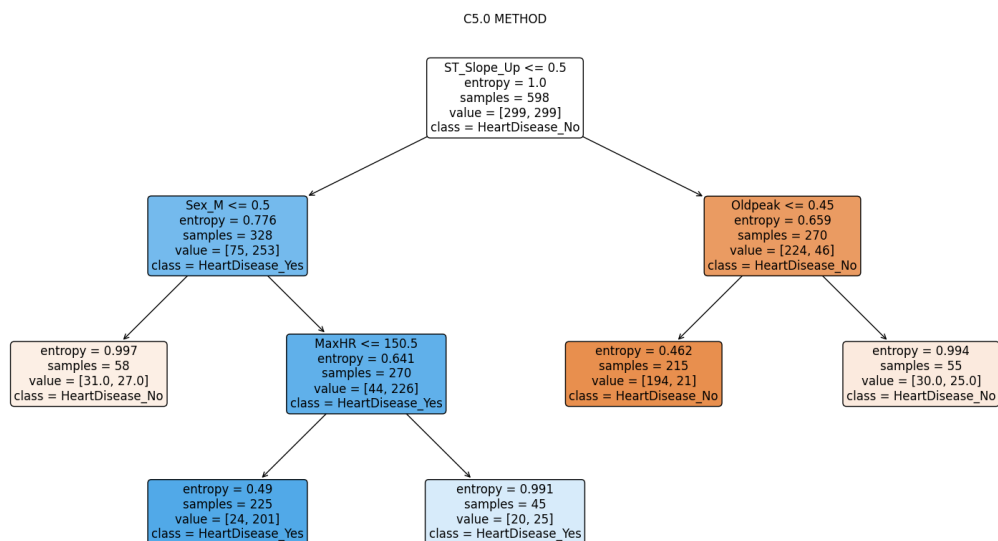
Accuracy Score:

0.8

C5.0 Method

```
[ ]: # Training algorithm
c50_01 = DecisionTreeClassifier(criterion= 'entropy', max_leaf_nodes= 5).
    ↪ fit(X,y)

# Visualizing c5.0
plt.figure(figsize=(20,10))
plot_tree(c50_01, feature_names=X_names, class_names=y_names, filled= True,
    ↪ rounded= True, fontsize= 12)
plt.title("C5.0 METHOD")
plt.show()
```



Predictions using c5.0

```
[ ]: predHeartDiseasec50 = c50_01.predict(X_test)

# Results
print("Confusion Matrix:")
print(confusion_matrix(y_test, predHeartDiseasec50))
print("\nClassification Report:")
print(classification_report(y_test, predHeartDiseasec50, target_names=["0",
↪ "1"]))
print("\nAccuracy Score:")
print(accuracy_score(y_test, predHeartDiseasec50))
```

Confusion Matrix:

```
[[100  11]
 [ 29  90]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.78	0.90	0.83	111
1	0.89	0.76	0.82	119
accuracy			0.83	230
macro avg	0.83	0.83	0.83	230
weighted avg	0.84	0.83	0.83	230

Accuracy Score:

0.8260869565217391

Using Random Forest Method

```
[ ]: # Requires a response variable formatted as a one-dimensional array
rfy = np.ravel(y)

# Training RandomForest
rf01 = RandomForestClassifier(n_estimators= 100, random_state=42, criterion=
↪ 'gini').fit(X,rfy)

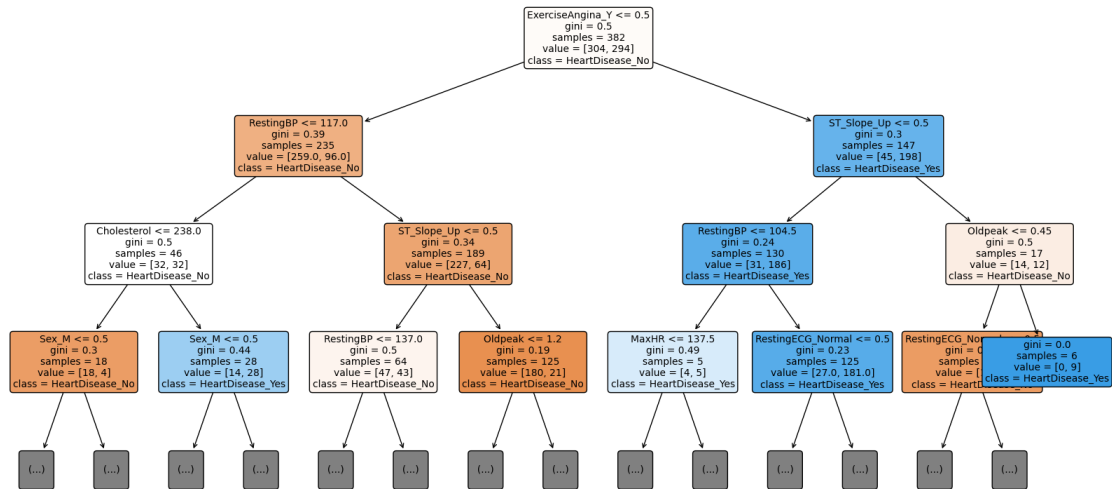
# Visualizing Random Forests
estimator = rf01.estimators_[0]

plt.figure(figsize=(20, 10))
tree.plot_tree(estimator,
                feature_names=X.columns,
                class_names=["HeartDisease_No", 'HeartDisease_Yes'],
                filled=True,
```

```

rounded=True,
proportion=False,
precision=2,
fontsize=10,
max_depth=3)
plt.show()

```



Predicting using RandomForest method

```

[ ]: predHeartDiseaseRandomForest = rf01.predict(X_test)

# Results
print("Confusion Matrix:")
print(confusion_matrix(y_test, predHeartDiseaseRandomForest))
print("\nClassification Report:")
print(classification_report(y_test, predHeartDiseaseRandomForest,
    target_names=["0", "1"]))
print("\nAccuracy Score:")
print(accuracy_score(y_test, predHeartDiseaseRandomForest))

```

Confusion Matrix:

```

[[ 91  20]
 [  9 110]]

```

Classification Report:

	precision	recall	f1-score	support
0	0.91	0.82	0.86	111
1	0.85	0.92	0.88	119

accuracy			0.87	230
macro avg	0.88	0.87	0.87	230
weighted avg	0.88	0.87	0.87	230

Accuracy Score:
0.8739130434782608

Naïve Bayes Method First, handle negative numbers in column “Oldpeak”

```
[ ]: # Initializing the MinMaxScaler function
scaler = MinMaxScaler()

X['Oldpeak'] = scaler.fit_transform(X[['Oldpeak']])
```


Training Naïve Bayes

```
[ ]: # Training Naive Bayes model
nb_01 = MultinomialNB().fit(X,rfy)

# Get feature log probabilities
feature_log_probs = nb_01.feature_log_prob_

# Convert to DataFrame for easier plotting
feature_log_probs_df = pd.DataFrame(feature_log_probs, columns=X.columns,
    ↪index=['Class0', 'Class1'])

# Plot the feature log probabilities as a heatmap
plt.figure(figsize=(12, 6))
sns.heatmap(feature_log_probs_df.T, annot=True, cmap='coolwarm', cbar=True,
    ↪linewidths=0.5)
plt.title('Feature Log Probabilities for Each Class')
plt.xlabel('Class')
plt.ylabel('Features')
plt.show()
```



Predicting using Naïve Bayes method

```
[ ]: predHeartDiseasecNB = nb_01.predict(X_test)

# Results
print("Confusion Matrix:")
print(confusion_matrix(y_test, predHeartDiseasecNB))
print("\nClassification Report:")
print(classification_report(y_test, predHeartDiseasecNB, target_names=["0", "1"]))
print("\nAccuracy Score:")
print(accuracy_score(y_test, predHeartDiseasecNB))
```

Confusion Matrix:

```
[[88 23]
 [20 99]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.81	0.79	0.80	111
1	0.81	0.83	0.82	119
accuracy			0.81	230
macro avg	0.81	0.81	0.81	230
weighted avg	0.81	0.81	0.81	230

Accuracy Score:
0.8130434782608695

Showing the contingency table of actual versus predicted outcomes

```
[ ]: ypred = pd.crosstab(y_test, predHeartDiseasecNB, rownames= ['Actual'],  
    ↪ colnames= ['Predicted'])  
  
ypred['Total'] = ypred.sum(axis=1)  
ypred.loc['Total'] = ypred.sum()  
ypred
```

```
[ ]: Predicted  False  True  Total  
Actual  
0             88    23    111  
1             20    99    119  
Total         108   122    230
```

88 cases: The actual value is 0, and the model correctly predicted 0 (True Negative).

23 cases: The actual value is 0, but the model incorrectly predicted 1 (False Positive).

20 cases: The actual value is 1, but the model incorrectly predicted 0 (False Negative).

99 cases: The actual value is 1, and the model correctly predicted 1 (True Positive).

Neural Networks

```
[ ]: X_train = X  
  
# Performing min-max standarization on all numerical variables  
X_train.loc[:, numerical_columns] = scaler.  
    ↪ fit_transform(X_train[numerical_columns])  
  
# Building keras model  
model = Sequential()  
model.add(Input(shape=(15,)))  
model.add(Dense(10, activation='relu'))  
model.add(Dense(1, activation='sigmoid'))  
  
# Compile keras model  
model.compile(optimizer='adam', loss='binary_crossentropy',  
    ↪ metrics=['accuracy'])  
  
# Train the model  
model.fit(X_train, y, epochs=150, verbose=0)  
  
# evaluate the keras model  
_, accuracy = model.evaluate(X_train, y, verbose=0)  
print('Accuracy: %.2f' % (accuracy*100))
```

Accuracy: 85.79

Making predictions

```
[ ]: # Performing min-max standarization on all numerical variables
X_test.loc[:, numerical_columns] = scaler.transform(X_test[numerical_columns])

# Predictions
prednnet01 = model.predict(X_test)
```

8/8 0s 857us/step

8/8 0s 857us/step

```
[ ]: # Convert probabilities to binary predictions (0 or 1)
threshold = 0.5
prednnet01_binary = (prednnet01 > threshold).astype(int)

# Reshape to match y_test if necessary
prednnet01_binary = prednnet01_binary.reshape(-1)

# Print confusion matrix
print("Confusion Matrix:")
print(confusion_matrix(y_test, prednnet01_binary))

# Print classification report
print("\nClassification Report:")
print(classification_report(y_test, prednnet01_binary, target_names=['No Heart_
↵Disease', 'Heart Disease']))

# Print accuracy score
print("\nAccuracy Score:")
print(accuracy_score(y_test, prednnet01_binary))
```

Confusion Matrix:

```
[[ 83  28]
 [ 10 109]]
```

Classification Report:

	precision	recall	f1-score	support
No Heart Disease	0.89	0.75	0.81	111
Heart Disease	0.80	0.92	0.85	119
accuracy			0.83	230
macro avg	0.84	0.83	0.83	230
weighted avg	0.84	0.83	0.83	230

Accuracy Score:

0.8347826086956521

Conclusion CART model gives the accuracy of : 80%

C5.0 model gives the accuracy of : 83%

Random forest gives the accuracy of : 86%

Naïve Bayes gives the accuracy of : 81%

Neural Network gives the accuracy of : 87%